



Stock Forecasting with Machine Learning

Caila Hanus, Paulette Petracco, Dani Bar-Lavi,
Kayode Ayenioye



Background

The stock market's dynamic and volatile nature presents both significant opportunities and considerable risks. Accurately forecasting stock prices can give investors a competitive edge, helping them make informed decisions and optimize their portfolios. Our project aims to leverage machine learning techniques to enhance the accuracy of stock price predictions and assess the associated risks.

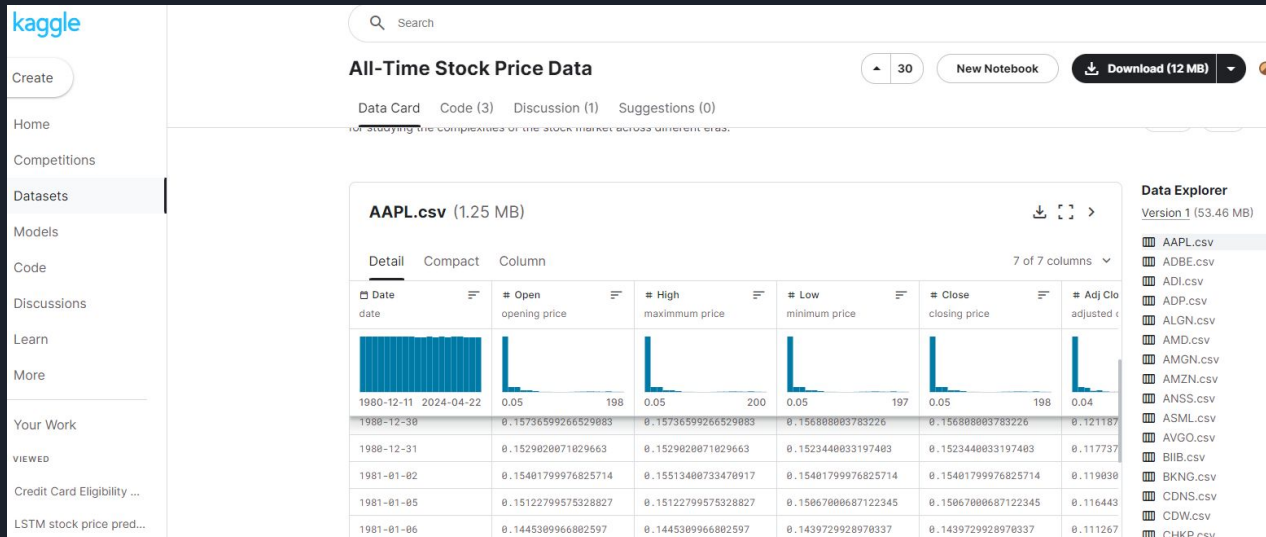
Main Research Questions:

- Predicting Stock Market Movement:
 - Short-Term Predictions: Where will the stock be tomorrow? Next month?
 - Long-Term Predictions: Where will the stock be next year? Next five years?
- Sentiment Analysis:
 - Financial News Impact: Analyze the sentiment of financial news and its impact on stock prices.
 - Social Media Influence: Examine how social media sentiment affects stock price predictions.

Data (Extraction)

Dataset: <https://www.kaggle.com/datasets/hchsmost/test-dataset>

- 77 .csv files including historical stock price data spanning various time periods
- Records that span multiple decades to explore the dynamics of different stocks, industries, and market sectors




A decorative graphic on the left side of the slide. It consists of a blue parallelogram and a light green parallelogram, both tilted at an angle. The blue shape is in the foreground, and the green shape is partially behind it. The background of the entire slide is dark blue with subtle diagonal lines.

Short Term Forecasting

Model Troubleshooting - One Month Predictions

NFLX Dataset:

- 5517 rows of data ranging from 5/23/2002 to 4/23/2022 or 20 years of stock data
- Aimed to get predictions for the next month of data



NETFLIX

Model Optimization

Things experimented with:

- Data Preprocessing:
 - Handling missing values and Feature scaling/normalization
 - Ex. Experimenting with different cutoff dates.
 - Ex. Using training sets of specified dates.
- Feature Engineering:
 - Creation of lagged features and additional features
 - Ex. Monthly Averages, Day-to-Day Changes, Volatility.

```
# Load the data
data = pd.read_csv('Resources/NFLX.csv')

# Convert the 'Date' column to datetime format
data['Date'] = pd.to_datetime(data['Date'])

# Ensure the data is sorted by date
data = data.sort_values(by='Date')
✓ 0.0s

# Creating moving average features
data['MA_5'] = data['Close'].rolling(window=5).mean()
data['MA_10'] = data['Close'].rolling(window=10).mean()

# Drop the rows with NaN values (caused by rolling windows)
data = data.dropna()

# Define the target variable and features
features = ['Open', 'High', 'Low', 'Volume', 'MA_5', 'MA_10']
target = 'Close'

# Split the data into training and testing sets
train_data = data[data['Date'] < '2022-01-01']
test_data = data[data['Date'] >= '2022-01-01']

X_train = train_data[features]
y_train = train_data[target]
X_test = test_data[features]
y_test = test_data[target]
✓ 0.0s
```

Model Optimization

More Experimenting:

- Model Selection and Comparison:
 - Linear Regression vs. Random Forest Regression.
 - Performance metrics (e.g., MSE, R-squared).
- Epoch-based Training:
 - Using neural networks.
 - Impact of different architectures.
- Train Model
- Create Visualizations

```
# Initialize the model
model = RandomForestRegressor(n_estimators=100, random_state=42)
```

```
# Train the model
model.fit(X_train, y_train)
```

```
# Make predictions
predictions = model.predict(X_test)
```

```
# Evaluate the model
mse = mean_squared_error(y_test, predictions)
print(f'Mean Squared Error: {mse}')
r2_score = r2_score(y_test, predictions)
print(f'R squared: {r2_score}')
```

```
✓ 9.5s
Mean Squared Error: 25.547305957802443
R squared: 0.9982296463424087
```

```
# Predict future prices for the next month
def predict_next_month(model, last_data, features, days=30):
    future_predictions = []
    last_date = last_data['Date'].max()

    for i in range(days):
        last_row = last_data.iloc[-1]
        new_row = last_row.copy()

        # Move the date forward by one day
        new_row['Date'] = last_date + pd.Timedelta(days=1)

        # Predict the next closing price
        prediction = model.predict([last_row[features]])[0]
        new_row['Close'] = prediction

        # Update the moving averages
        last_data = pd.concat([last_data, new_row.to_frame().T], ignore_index=True)
        last_data['MA_5'] = last_data['Close'].rolling(window=5).mean()
        last_data['MA_10'] = last_data['Close'].rolling(window=10).mean()

        # Append the prediction to the future predictions list
        future_predictions.append(new_row)

        # Update the last date
        last_date = new_row['Date']





        # Remove NaN values that may appear during the moving average calculation
        last_data = last_data.dropna()

    return pd.DataFrame(future_predictions)


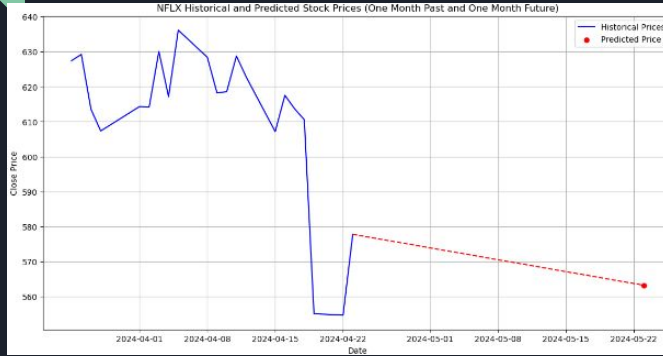
# Prepare the last data point for prediction
last_data = data.copy()

# Predict future prices
future_predictions = predict_next_month(model, last_data, features)
```

One Month Forecasting by Day



One Month Comparison to Forecasted





Predicting Stock Prices in the Health and Beauty Sector

The stocks analyzed in the 5 year model are:

- Align Technology, Inc. (ALGN)
- Lululemon Athletica Inc. (LULU)
- Peloton Interactive, Inc. (PTON)
- Ulta Beauty, Inc. (ULTA)

| | | | | |
|------|-------------------------|----------|----------|----------|
| ALGN | Align Technology, Inc. | \$255.70 | +\$0.16 | ↑ 0.063% |
| LULU | Lululemon Athletica Inc | \$318.26 | +\$0.40 | ↑ 0.13% |
| PTON | Peloton Interactive Inc | \$3.64 | -\$0.020 | ↓ 0.55% |
| ULTA | Ulta Beauty Inc | \$382.50 | -\$0.11 | ↓ 0.029% |

By leveraging 10 years of historical data, feature engineering, and neural networks, I attempted to forecast stock prices for the next five years.



Data Loading and Preparation

- Created a Library
- Pulled in the 77 files
- Combined the files

```
# Define directory containing the csv files
dir_path = Path("../Resources")

# create an empty list to hold the dataframes
dfs = []

# loop through each file in the directory
for file in os.listdir(dir_path):
    # check if the file is a csv file
    if file.endswith(".csv"):
        # extract the ticker symbol from the file name (assuming the file name is the ticker symbol)
        ticker = file.replace(".csv", "")
        # Read the file into a DataFrame
        stocks_df = pd.read_csv(dir_path / file)
        # Add a column to the DataFrame to store the ticker symbol
        stocks_df["Ticker"] = ticker
        # add the dataframe to the list
        dfs.append(stocks_df)

# concatenate the dataframes in the list
combined_stocks_df = pd.concat(dfs, ignore_index=True)

# Display the combined DataFrame to verify
print(combined_stocks_df.head())

# change the type in 'date' column to datetime
combined_stocks_df["Date"] = pd.to_datetime(combined_stocks_df["Date"])
```

Engineering and Model Training

```
stocks_to_analyze = ['ALGN', 'LULU', 'PTON', 'ULTA']

# Initialize dictionaries to hold the MSE and predictions for each stock symbol
mse_dict = {}
predictions_dict = {}

# Iterate through each stock in the list of specific stocks
for stock in stocks_to_analyze:
    # Select the stock's data
    stock_data = stocks_data_filtered[stocks_data_filtered["Ticker"] == stock]

    # Display the selected stock data
    print(f"Selected stock data for {stock}:")
    print(stock_data.head())

    # Set the date as the index
    stock_data.set_index("Date", inplace=True)

    # Define a feature in the data for previous date closing prices
    stock_data["Previous Day Close"] = stock_data["Close"].shift(1)

    # Define a feature in the data for the volume
    stock_data["Volume Difference"] = stock_data["Volume"].diff()

    # Drop rows with NaN values
    stock_data = stock_data.dropna()
```

Engineering

```
# Define the features (X) and the target (y) variables for training purposes
X = stock_data[["Previous Day Close", "Volume Difference"]]
y = stock_data["Close"].values.reshape(-1, 1)
```

```
# Split the data into training and testing sets chronologically
split = int(0.7 * len(X))
X_train = X[:split]
X_test = X[split:]
y_train = y[:split]
y_test = y[split:]
```

Model Training

```
# Normalize the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
# Define the neural network model
model = Sequential()
model.add(Dense(64, activation='relu', input_dim=X_train_scaled.shape[1]))
model.add(Dense(32, activation='relu'))
model.add(Dense(1))
```

```
# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')
```

```
# Train the model with epochs
model.fit(X_train_scaled, y_train, epochs=50, batch_size=32, validation_split=0.2)
```

```
# Make predictions using the testing data
predictions = model.predict(X_test_scaled)
```

Insights

Future predictions

Mean Squared Error

Stock Ticker, Mean Squared Error

ALGN, 138.9862474570898

LULU, 80.46665080758463

PTON, 248.35520640453953

ULTA, 113.22467256840763

Date, Close

2024-04-30, 312.58243

2024-05-31, 313.27362

2024-06-30, 313.95377

2024-07-31, 314.62308

2024-08-31, 315.2817

2024-09-30, 315.92987

2024-10-31, 316.5677

2024-11-30, 317.19534

2024-12-31, 317.813

2025-01-31, 318.42078

2025-02-28, 319.01892

2025-03-31, 319.6075

2025-04-30, 320.1867

2025-05-31, 320.75668

2025-06-30, 321.31754

2025-07-31, 321.8695

2025-08-31, 322.4126

2025-09-30, 322.94708

2025-10-31, 323.47305

2025-11-30, 323.99057

2025-12-31, 324.49988

ALGN

2026-01-31, 325.0011

2026-02-28, 325.4943

2026-03-31, 325.9796

2026-04-30, 326.4572

2026-05-31, 326.9272

2026-06-30, 327.38968

2026-07-31, 327.8448

2026-08-31, 328.2926

2026-09-30, 328.7333

2026-10-31, 329.16702

2026-11-30, 329.5938

2026-12-31, 330.01382

2027-01-31, 330.42706

2027-02-28, 330.83374

2027-03-31, 331.23395

2027-04-30, 331.62772

2027-05-31, 332.01526

2027-06-30, 332.39664

2027-07-31, 332.7719

2027-08-31, 333.1412

2027-09-30, 333.50455

2027-10-31, 333.86215

2027-11-30, 334.21405

2027-12-31, 334.56036

2028-01-31, 334.90115

2028-02-29, 335.23648

2028-03-31, 335.56647

2028-04-30, 335.89117

2028-05-31, 336.21072

2028-06-30, 336.52518

2028-07-31, 336.83463

2028-08-31, 337.13916

2028-09-30, 337.4388

2028-10-31, 337.7337

2028-11-30, 338.0239

2028-12-31, 338.30945

2029-01-31, 338.59048

2029-02-28, 338.86697

2029-03-31, 339.1391

```
# Predict the future values iteratively
future_closes = []
for date in future_dates_df.index:
    # prepare the input data for prediction
    input_data = pd.DataFrame({"Previous Day Close": [previous_close], "Volume Difference": [volume_difference]})
    input_data_scaled = scaler.transform(input_data)
    # Make the prediction
    predicted_close = model.predict(input_data_scaled)[0][0]
    # Append the predicted close to the list
    future_closes.append(predicted_close)
    # Update previous_close for the next iteration
    previous_close = predicted_close

# Add the predictions to the dataframe
future_dates_df["Close"] = future_closes

# Store the predictions in the dictionary
predictions_dict[stock] = future_dates_df
```

```
# Display the predictions for each stock ticker
for stock, future_dates_df in predictions_dict.items():
```

```
    print(f"Predictions for {stock}:")
    print(future_dates_df)

    # Define the base path as the parent directory of "Notebooks" and "Resources"
    base_path = os.path.abspath(os.path.join(os.getcwd(), os.pardir))

    # Create a directory named "Predictions" within the base path if it doesn't exist
    predictions_folder = os.path.join(base_path, "Health and Beauty Five Year Predictions")
    if not os.path.exists(predictions_folder):
        os.makedirs(predictions_folder)

    # Save the Predictions results to individual files in the "Predictions" folder
    for stock, future_dates_df in predictions_dict.items():
        prediction_file_path = os.path.join(predictions_folder, f"predictions_{stock}.csv")
        future_dates_df.to_csv(prediction_file_path)

    print(f"Prediction results saved to the folder: {predictions_folder}")


    # Create a directory named "MSE" within the base path if it doesn't exist
    mse_folder = os.path.join(base_path, "Health and Beauty Five Year MSE Output")
    if not os.path.exists(mse_folder):
        os.makedirs(mse_folder)

    # Define the path to save the MSE results file
    mse_file_path = os.path.join(mse_folder, "mse_results.csv")

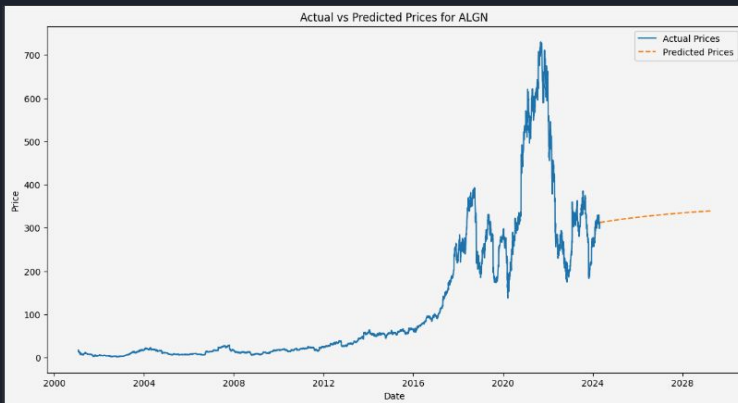
    # Save the MSE results to a file in the "MSE" folder
    mse_df = pd.DataFrame(list(mse_dict.items()), columns=["Stock Ticker", "Mean Squared Error"])
    mse_df.to_csv(mse_file_path, index=False)

    print(f"MSE results saved to: {mse_file_path}")
```

Reviewing The Last Ten Years of Data to Predict The Next Five



5 Year - Actual vs Predicted





Modeling Big Stock Data

- Wrote one code to iterate a model over the historical data of 77 stocks
- Using nesting for-loops, the model is able to run separately over each stock and give separate precision and accuracy scores for the model as it pertains to each stock
- Created three slightly different versions of my code to create 6 month, 12 month, and five year projections. I'll be focusing on 12 month projections for this presentation.



Process

1. Data Loading and Preprocessing

- a. Filter data by date
- b. Feature engineering—create new features for previous days close and volume

```
# Use a for loop to iterate data_load over each CSV in the Data folder to get individual results for each stock
for stock in os.listdir(data_path):
    if stock.endswith('.csv'):
        file_path = os.path.join(data_path, stock)
        data = data_load(file_path)

        # Filter this stock's data to only use entries from January 1st, 2020 onward
        start_date = datetime(2020, 1, 1)
        data_filtered_2020s = data[data['Date'] >= start_date].copy()

        # Define new features in the data for the previous day's close and previous day's volume
        data_filtered_2020s.loc[:, 'Prev_Day_Close'] = data_filtered_2020s['Close'].shift(1)
        data_filtered_2020s.loc[:, 'Prev_Day_Vol'] = data_filtered_2020s['Volume'].shift(1)

        # Drop rows with NaN values using dropna()
        data_filtered_2020s = data_filtered_2020s.dropna()

        # Ensure data is sorted by date
        data_filtered_2020s = data_filtered_2020s.sort_values(by='Date')
```



Process

2. Preparing Data for Testing:

- Feature/Target Definition
- Data Splitting for Testing (Chronological)

```
# Define features (X) and target (y) for training purposes  
X = data_filtered_2020s[['Prev_Day_Close', 'Prev_Day_Vol']]  
y = data_filtered_2020s['Close']
```

```
# Split data into test and training sets chronologically  
cutoff_date = datetime(2023, 12, 31)  
X_train = X[data_filtered_2020s['Date'] <= cutoff_date]  
y_train = y[data_filtered_2020s['Date'] <= cutoff_date]  
X_test = X[data_filtered_2020s['Date'] > cutoff_date]  
y_test = y[data_filtered_2020s['Date'] > cutoff_date]
```



Process

3. Model Training

4. Model Evaluation

```
# Train stock_model  
stock_model = RandomForestRegressor(n_estimators=100, random_state=42)  
stock_model.fit(X_train, y_train)
```

```
# Evaluate the stock_model using mse (mean squared error)  
y_pred = stock_model.predict(X_test)  
mse = mean_squared_error(y_test, y_pred)  
print(f'{stock}: Mean Squared Error: {mse}')
```

Process

5. Make Future Predictions (Iteratively)

```
# Make full year predictions for this stock
future_dates = pd.date_range(start='2024-07-01', end='2025-7-01', freq='B') # 'B' restricts to business days
future_data = pd.DataFrame({
    'Date': future_dates
})

# Initialize with the last known close
last_known_close = data_filtered_2020s['close'].iloc[-1]

# Define the range for random volume generation
min_vol = min(data_filtered_2020s['Volume'])
max_vol = max(data_filtered_2020s['Volume'])

# Iteratively predict each day's closing price using the previous day's projected close
projected_closes = []
for i in range(len(future_dates)):
    if i == 0:
        prev_close = last_known_close
    else:
        prev_close = projected_closes[-1]

    prev_vol = np.random.uniform(low=min_vol, high=max_vol)

    future_data.loc[i, 'Prev_Day_Close'] = prev_close
    future_data.loc[i, 'Prev_Day_Vol'] = prev_vol

    # Predict the closing price
    pred_close = stock_model.predict([[prev_close, prev_vol]])[0]
    projected_closes.append(pred_close)

future_data['Projected_Close'] = projected_closes
future_data['Stock'] = stock.replace('.csv', '') # Add stock identifier
stock_predictions.append(future_data)
```



Process

6. Concatenating and saving projection data


```
# Concatenate all predictions into one DataFrame
stock_predictions_df = pd.concat(stock_predictions, ignore_index=True)

print(stock_predictions_df.head())


# Save predictions to a CSV file
stock_predictions_df.to_csv('fullyr_future_stock_predictions.csv', index=False)
print('Stock projections exported to fullyr_future_stock_predictions.csv')
```

Visualizations


These are the projections for the top 3 lowest MSE stocks I analyzed, starting with Sirius XM.



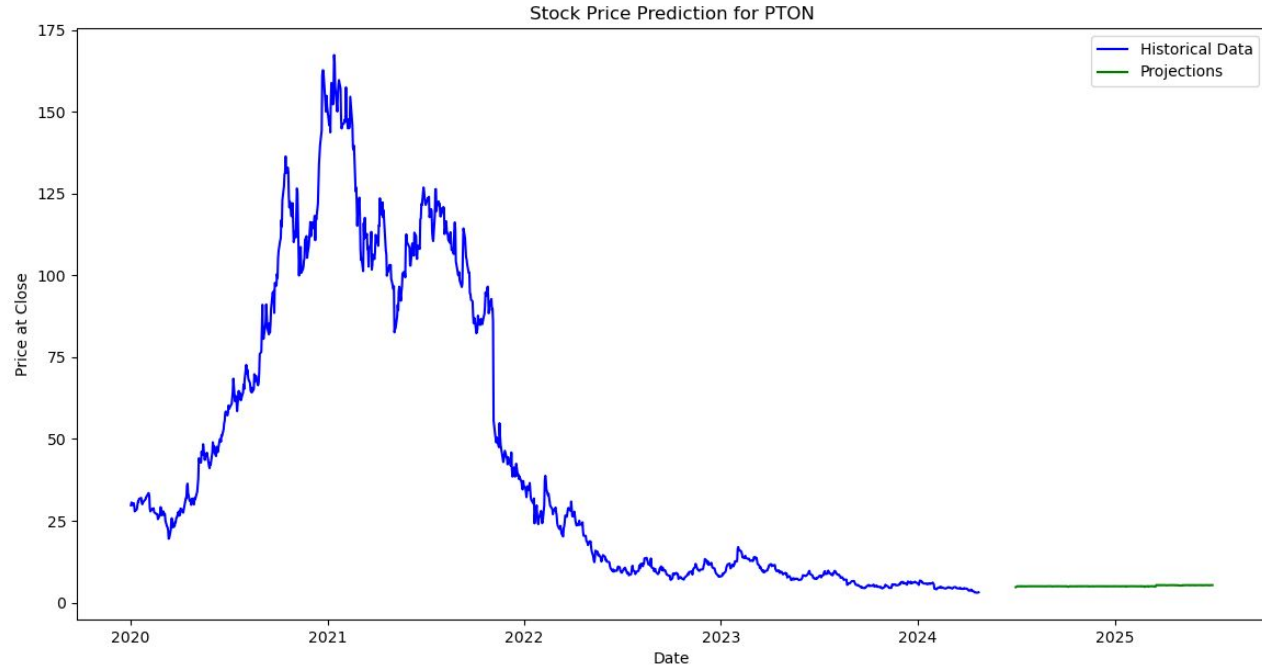
20th Century Fox



Kraft Heinz



Peloton and Data Storytelling




Reliance Stock Price Prediction

Historical Stock Prices

Volatility Assessment

Sentiment Analysis

5-Year Stock Price Projection





Why Sentiment Analysis of Financial News and Social Media?

- Analyzing market sentiment from news articles, social media, etc
- Creating market sentiment indicators indicating bullish or bearish trends
- Understanding how positive or negative news influences investor behaviour
- Monitoring sentiments on social media platforms to gauge public perception and sentiment towards stocks or the market